

Topic 5 — Random generation, I

5.1 Simulating simple discrete distributions with a fair coin

The simplest of all distributions is the fair coin. Let's code its two outcomes, heads and tails, by 1 and 0 respectively.

0 with probability $1/2$
 1 with probability $1/2$

We call this the Bernoulli($1/2$) or $B(1/2)$ distribution.

With just the ability to flip fair coins (or equivalently, to sample from the $B(1/2)$ distribution), we can generate random samples from any discrete distribution with a finite sample space. We will get to this result in several steps.

5.1.1 Uniform distribution over b -bit integers

The uniform distribution over b -bit integers has the following probability space:

$$\begin{aligned}\Omega &= \{0, 1\}^b \\ \Pr(\omega) &= 1/2^b \text{ for all } \omega \in \Omega\end{aligned}$$

Call this distribution $\text{Unif}(\{0, 1\}^b)$. To sample from it, simply flip a fair coin for each of the b bits.

```
For  $i = 1$  to  $b$ :
    Draw  $X_i$  from  $B(1/2)$ 
Output  $X_1 X_2 \dots X_b$ 
```

5.1.2 Uniform distribution over $\{1, 2, \dots, n\}$

Now consider a very similar distribution with probability space

$$\begin{aligned}\Omega &= \{1, 2, 3, \dots, n\} \\ \Pr(\omega) &= 1/n \text{ for all } \omega \in \Omega\end{aligned}$$

Call this distribution $\text{Unif}(\{1, \dots, n\})$. If n is of the form 2^b , then the previous algorithm, called with $b = \log_2 n$, gives us a uniform distribution over $\{0, 1, \dots, n-1\}$. So we can just add 1 to that value and we're done.

More generally, here's a sampling algorithm.

```
Let  $b = \lceil \log_2 n \rceil$ 
Repeat:
    Generate a sample  $X$  from  $\text{Unif}(\{1, 2, \dots, 2^b\})$ , as described above
    If  $X \leq n$ : output  $X$  and halt
```

First, let's check that this indeed outputs the right distribution, that each of the values $1, 2, \dots, n$ gets output with probability exactly $1/n$. Specifically, we need to show that if X is a sample from $\text{Unif}(\{1, 2, \dots, 2^b\})$, then for any $i \in \{1, 2, \dots, n\}$, we have $\Pr(X = i | X \leq n) = 1/n$. This follows from the formula for conditional probability:

$$\Pr(X = i | X \leq n) = \frac{\Pr(X = i \text{ AND } X \leq n)}{\Pr(X \leq n)} = \frac{\Pr(X = i)}{\Pr(X \leq n)} = \frac{1/2^b}{n/2^b} = \frac{1}{n}.$$

How many coin flips does this algorithm use? Each time we go through the repeat loop, we use b flips to generate X ; but how many times do we loop? First notice that $b = \lceil \log_2 n \rceil$, which means that b is the smallest integer that is greater than or equal to $\log_2 n$. Therefore

$$b - 1 < \log_2 n \leq b \Rightarrow \frac{1}{2} 2^b < n \leq 2^b \Rightarrow \Pr(X \leq n) = \frac{n}{2^b} > \frac{1}{2}.$$

Therefore, on each iteration of the repeat loop, the probability of halting, $\Pr(X \leq n)$, is at least $1/2$. So the expected number of iterations is at most 2, which means that the expected number of coin flips needed is at most $2b$.

5.1.3 Uniform distribution over $[0, 1]$

This time, we want a uniform distribution over real numbers in the interval $[0, 1]$. However, the size of this sample space is uncountably infinite, and for a variety of practical reasons, we will typically want only a finite amount of precision.

Recall the binary representation of fractional values: $0.z_1z_2z_3\dots$. Here z_1 is the position for $1/2$, z_2 is the position for $1/4$, z_3 is the position for $1/8$, and so on. For instance, $0.101 = 1/2 + 1/8 = 5/8$ whereas $0.0011 = 1/8 + 1/16 = 3/16$.

Let's say that we want b bits of precision.

$$\begin{aligned}\Omega &= \{0.z_1z_2\dots z_b : z_1, \dots, z_b \in \{0, 1\}\} \\ \Pr(\omega) &= 1/2^b \text{ for all } \omega \in \Omega\end{aligned}$$

This is exactly like generating a random b -bit integer: just stick a "0." in front.

5.1.4 A biased coin

The next distribution we want is a coin with bias p , where the outcome is once again coded as 0/1:

$$\begin{aligned}0 &\text{ with probability } 1 - p \\ 1 &\text{ with probability } p\end{aligned}$$

We call this the Bernoulli(p) or $B(p)$ distribution. Can we simulate $B(p)$ using $B(1/2)$?

Here's an easy way to do so.

```
Generate  $X$  from  $\text{Unif}[0, 1]$ 
If  $X \leq p$ : output 1
else:      output 0
```

Since $\Pr(X \leq p) = p$, this generates the right distribution. But how many fair coin flips does it use? As stated here, it seems to require that X is infinite precision. The way around this is to notice that we can just generate X one bit at a time, until it is clear whether X is less than p or more than p .

For instance, suppose $p = 3/8$. In binary, this is 0.011. Writing $X = 0.X_1X_2X_3\cdots$, we first flip a coin to get X_1 , then another coin to get X_2 , and so on. Suppose $X_1 = 1$. Then we can stop at once, because we know that X is at least $1/2$ and therefore $X \geq p$, no matter what X_2, X_3, \dots turn out to be. On the other hand, if $X_1 = 0$, then all we know is that $X \leq 1/2$, so we can't be sure whether it is bigger or smaller than p , and we have to continue. Here's the modified algorithm:

```

Let  $0.p_1p_2p_3\cdots$  be the binary representation of  $p$ 
Repeat for  $i = 1, 2, 3, \dots$  :
    Draw  $X_i$  from  $B(1/2)$ 
    If  $p_i = 1$  and  $X_i = 0$ :  halt and output 1
    If  $p_i = 0$  and  $X_i = 1$ :  halt and output 0

```

How many bits are needed? That is, how many times does the algorithm loop? Notice that on each iteration, the algorithm halts if $X_i \neq p_i$. This happens with probability exactly $1/2$. Therefore, the expected number of iterations is exactly 2.

So we can simulate a biased coin using, on average, two fair coins.

5.1.5 Arbitrary discrete distribution with finite sample space

Let's move to a much more general distribution.

$$\begin{aligned}\Omega &= \{\omega_1, \dots, \omega_k\} \\ \Pr(\omega_i) &= p_i\end{aligned}$$

where the p_i are nonnegative and sum to 1. An example is the roll of a die, which has $k = 6$ and $p_1 = \dots = p_k = 1/6$.

To sample from this distribution, we use the same ideas as for a biased coin. Let's start with the infinite precision version.

```

Generate  $X$  from  $\text{Unif}[0, 1]$ 
For all  $i = 1$  to  $k$ :
    If  $p_1 + \dots + p_{i-1} < X \leq p_1 + \dots + p_i$ :  output  $\omega_i$ 

```

In effect, we divide the interval $[0, 1]$ into k bins, where the i th bin stretches from $p_1 + \dots + p_{i-1}$ to $p_1 + \dots + p_i$, and therefore has length exactly p_i . We generate X uniformly from $[0, 1]$ and then output the index of the bin that it falls into. The chance of falling into the i th bin (that is, of outputting ω_i) is therefore exactly p_i .

As before, we can run this process by generating X one bit at a time, and stopping as soon as it is clear which bin X will fall into. It is possible to show that the expected number of bits (coin flips) needed is at most $1 + \log_2 k$.