

Machine learning versus Algorithms

In both fields, the goal is to develop

procedures that exhibit a desired input-output behavior.

CSE 250B: Machine Learning

Lecture 0: Overview

- **Algorithms:** the input-output mapping can be precisely defined.
Input: Graph G .
Output: MST of G .
- **Machine learning:** the mapping cannot easily be made precise.
Input: Picture of an animal.
Output: Name of the animal.


Instead, we simply provide examples of (input,output) pairs and ask the machine to *learn* a suitable mapping itself.

Inputs and outputs

Basic terminology:

- The input space, \mathcal{X} .
E.g. 32×32 RGB images of animals.
- The output space, \mathcal{Y} .
E.g. Names of 100 animals.

x :



y : "bear"

After seeing a bunch of examples (x, y) , pick a mapping

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

that accurately replicates the input-output pattern of the examples.

Learning problems are often categorized according to the type of *output space*: (1) discrete, (2) continuous, (3) probability values, or (4) more general structures.

Discrete output space: classification

Binary classification:

- **Spam detection**
 $\mathcal{X} = \{\text{email messages}\}$
 $\mathcal{Y} = \{\text{spam, not spam}\}$
- **Credit card fraud detection**
 $\mathcal{X} = \{\text{descriptions of credit card transactions}\}$
 $\mathcal{Y} = \{\text{fraudulent, legitimate}\}$

Multiclass classification:

- **Animal recognition**
 $\mathcal{X} = \{\text{animal pictures}\}$
 $\mathcal{Y} = \{\text{dog, cat, giraffe, } \dots\}$
- **News article classification**
 $\mathcal{X} = \{\text{news articles}\}$
 $\mathcal{Y} = \{\text{politics, business, sports, } \dots\}$

Continuous output space: regression

- A parent's concerns

How cold will it be tomorrow morning?

$\mathcal{Y} = [-273, \infty)$

- For the asthmatic

Predict tomorrow's air quality (max over the whole day)

$\mathcal{Y} = [0, \infty)$ (< 100: okay, > 200: dangerous)

- Insurance company calculations

In how many years will this person die?

$\mathcal{Y} = [0, 200]$

What are suitable predictor variables (\mathcal{X}) in each case?

Conditional probability functions

Here $\mathcal{Y} = [0, 1]$ represents probabilities.

- Dating service

What is the probability these two people will go on a date if introduced to each other?

If we modeled this as a classification problem, the binary answer would basically always be “no”. The goal is to find matches that are slightly less unlikely than others.

- Credit card transactions

What is the probability that this transaction is fraudulent?

The probability is important, because – in combination with the amount of the transaction – it determines the overall risk and thus the right course of action.

Structured output spaces

The output space consists of structured objects, like sequences or trees.

Dating service

Input: description of a person

Output: rank-ordered list of all possible matches

\mathcal{Y} = space of all permutations

Example:

$x = \text{Tom}$

$y = (\text{Nancy}, \text{Mary}, \text{Chloe}, \dots)$

Language processing


Input: English sentence

Output: parse tree showing grammatical structure

\mathcal{Y} = space of all trees

Example:

$x = \text{“John hit the ball”}$

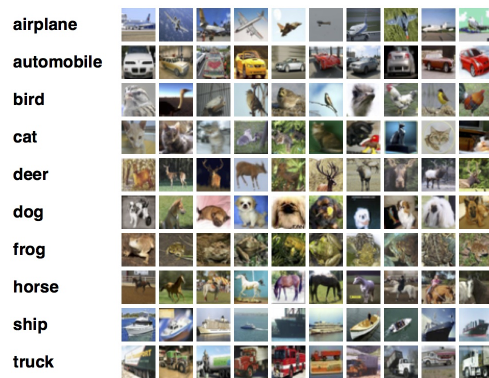


Course outline

- 1 Nonparametric methods
- 2 Classification using parametrized models
- 3 Combining classifiers
- 4 Representation learning

Nonparametric methods: nearest neighbor

Training set: a collection of (x, y) pairs:




Given any x , find its nearest neighbor in the training set and predict that neighbor's y value.

Issues: (1) What distance function? (2) How to speed up search?

Classification with parametrized models

Classifiers with a fixed number of parameters can represent a limited set of functions. Learning a model is about picking a good approximation.

Typically the x 's are points in d -dimensional Euclidean space, \mathbb{R}^d :

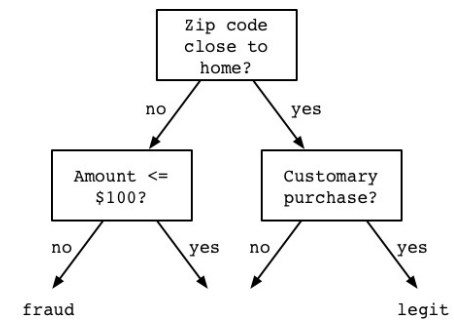


Two ways to classify:

- *Generative*: model the individual classes.
- *Discriminative*: model the decision boundary between the classes.

Nonparametric methods: decision tree

Credit card fraud detection: use training data to build a tree classifier



What do nearest neighbor and decision trees have in common?


- Unbounded in size
- Can model arbitrarily complex functions

They are *nonparametric methods*.

Generative models

Fit a probability distribution – like a multivariate Gaussian – to each class. Thereafter use this *summary* rather than the data points themselves.

To classify a new point: find the most probable class.




Examples: Naive Bayes, Fisher discriminant.

Under the hood: Bayes' rule, linear algebra (eigenvalues, eigenvectors).

Discriminative models

Approximate the boundaries between classes by simple – e.g. linear – functions.

To classify a new point: figure out which side of the boundary it lies on.



Examples: support vector machine, logistic regression.


Under the hood: convex duality, optimization.

Richer classifiers via the kernel trick

We are good at finding linear classifiers in Euclidean space. But what if:

- The boundary between classes is far from linear?

Example: quadratic, or higher-order polynomial, or even stranger.



- The data aren't even vectors of numbers?

Example: documents, DNA sequences, parse trees.

The *kernel trick* handles these scenarios seamlessly, by mapping the data to a suitable Euclidean space in which linear classification is possible!


Generalization theory

- Complex, e.g. nonparametric, classifiers require a lot of training data to learn accurately.
- Simple, e.g. linear, classifiers require less.

What is the right notion of complexity? Are there formulas for how much data is enough? The answers are based on *large deviation theory*.

Example: *Symmetric random walk*.

A drunken man starts at the origin and at each time step either takes a step to the right or a step to the left. Where is he after n steps?



Answer: Not far from where he started! Most likely in $[-c\sqrt{n}, c\sqrt{n}]$ (for some constant c), and all positions in this interval are about equally likely.

Under the hood: probability theory.

Richer output spaces

Many classification methods were developed for the binary (two-label) case. Usually the output space is larger than this.

- \mathcal{Y} = several classes.

Examples:

x = image, y = name of object in image

x = news article, y = category (sports, politics, business, ...)

- \mathcal{Y} = structured objects.

Examples:

x = sentence in Swahili, y = transcription into English

x = sentence in English, y = parse tree

Extend binary classification to handle such cases!

Under the hood: error-correcting codes, dynamic programming.


Composing simple classifiers

A common situation in classifier learning:

*Easy to find **weak classifiers** – not very accurate, but better than random*

To increase accuracy, *compose* weak classifiers.

Example: *boosting*.



Final classifier is a linear combination of all these weak classifiers.


Generically improve the performance of *any* kind of classifier!

Representation learning

A handful of key primitives:

- 1 Dimensionality reduction and denoising.
- 2 Embedding and manifold learning.

Given data that lie in a non-Euclidean space, find an embedding into Euclidean space that preserves as much of the geometry as possible.




- 3 Metric learning.

Representation learning

A handful of key primitives:

- 1 Dimensionality reduction and denoising.

Given data in high-dimensional Euclidean space, project to a low-dimensional linear subspace while retaining as much of the signal as possible.



- 2 Embedding and manifold learning.
- 3 Metric learning.

Representation learning

A handful of key primitives:

- 1 Dimensionality reduction and denoising.
- 2 Embedding and manifold learning.
- 3 Metric learning.

Given data with only vague positional information, impose an Euclidean geometry that is suitable for classification.

Example: $\mathcal{X} = \{\text{a collection of } m \text{ books}\}$.

A user supplies $\binom{m}{2}$ similarity ratings, such as:

("Pride and Prejudice", "Great Expectations"): similar


("Hamlet", "Great Expectations"): dissimilar

Represent each book by a vector, respecting these ratings.

Under the hood: linear algebra, semidefinite programming.

Deep learning

Multi-layer neural nets achieve state-of-the-art performance across a range of benchmark problems in natural language processing, speech, and vision.



Under the hood: stochastic gradient descent, dictionary learning, autoencoders.

Class details

Website: <http://www.cs.ucsd.edu/~dasgupta/250B>.

Grading:

- Regular homeworks: 50%.
- Two midterms: 25% each.

Course outline

- 1 Nonparametric methods
- 2 Classification using parametrized models
 - Generative models
 - Discriminative models
 - Richer decision boundaries using the kernel trick
 - Richer output spaces
 - Generalization theory
- 3 Combining classifiers
 - Boosting, bagging, and random forests
 - Online learning
- 4 Representation learning
 - Linear projection
 - Embeddings
 - Metric learning
 - Deep learning